# Distributed Dynamic Stream Processing
# Master Project WS 2020

Tilmann Rabl

Data Engineering Systems

# This Session

1. Intro Data Engineering Systems
2. Motivation Master Project
3. Milestones

# Data Engineering Systems

Systems that enable data engineering

- Database systems
- Stream processing systems
- Graph processing systems
- Machine learning systems
- …

Systems that support data science

- Experiment databases
- Optimizers
- Deployment systems
- End-to-end ML

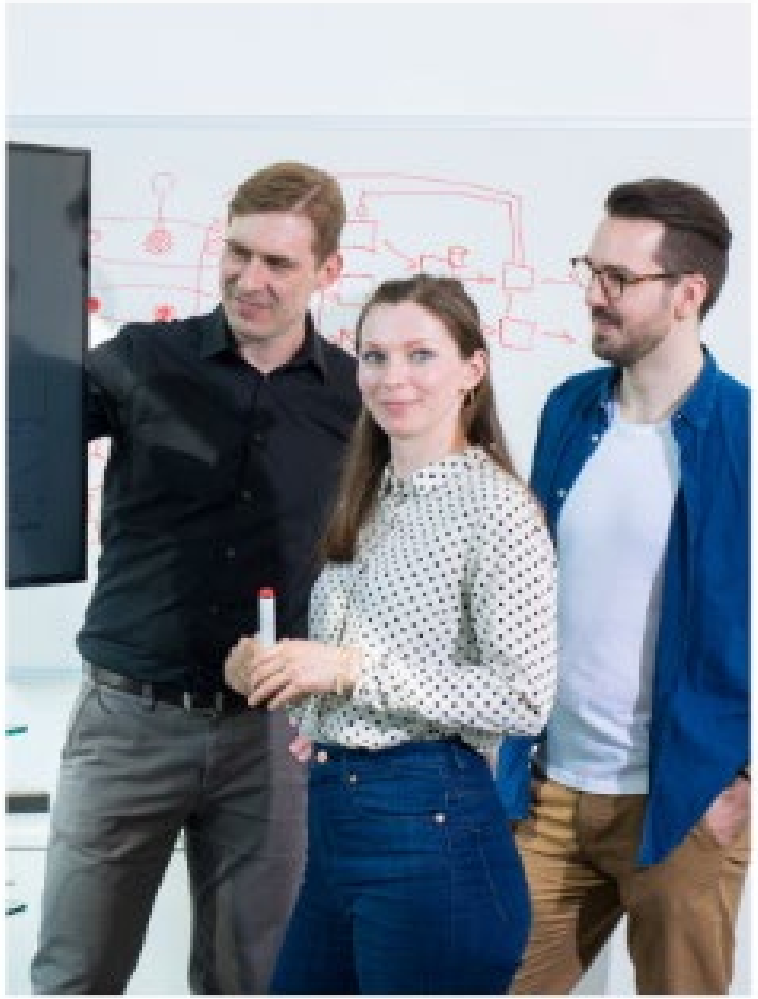

# Data Engineering Systems Group











# Research Topics

- Database Management
  - SIGMOD 17, VLDBJ 18, SIGMOD 20
- Machine Learning Systems
  - PVLDB 17,  SOCC 18, EDBT 20, SIGMOD 20
- Stream Processing
  - SIGMOD 19, PVLDB 19, SIGMOD 20, PVLDB 20
- Benchmarking
  - SOCC 17, ICDE 18, ICPE 18, PVDLB 20

Interested in a thesis? Write us an e-mail.

## Research Approach

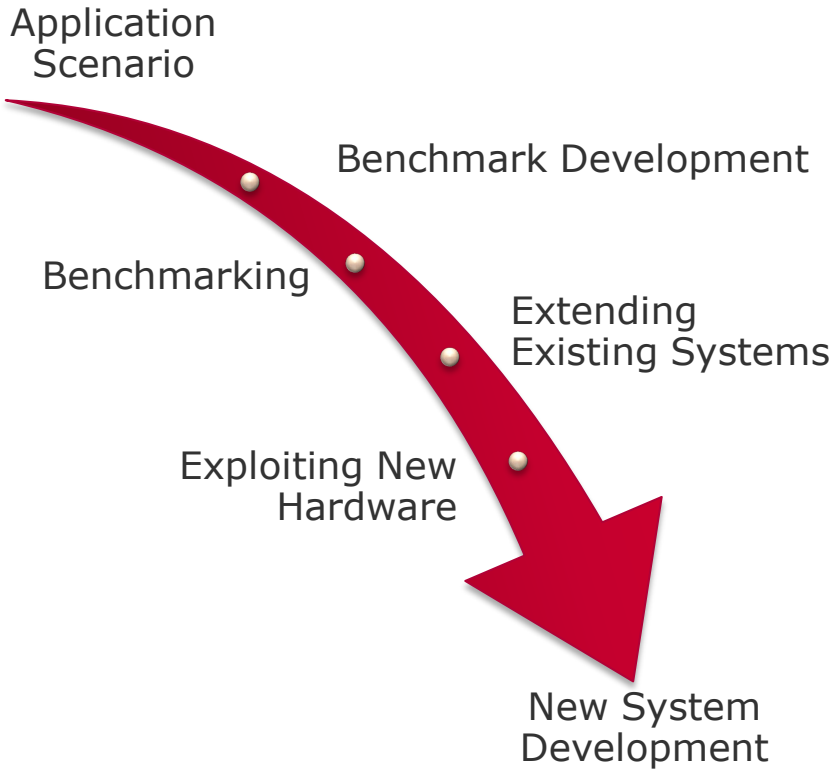

# Data Engineering Courses

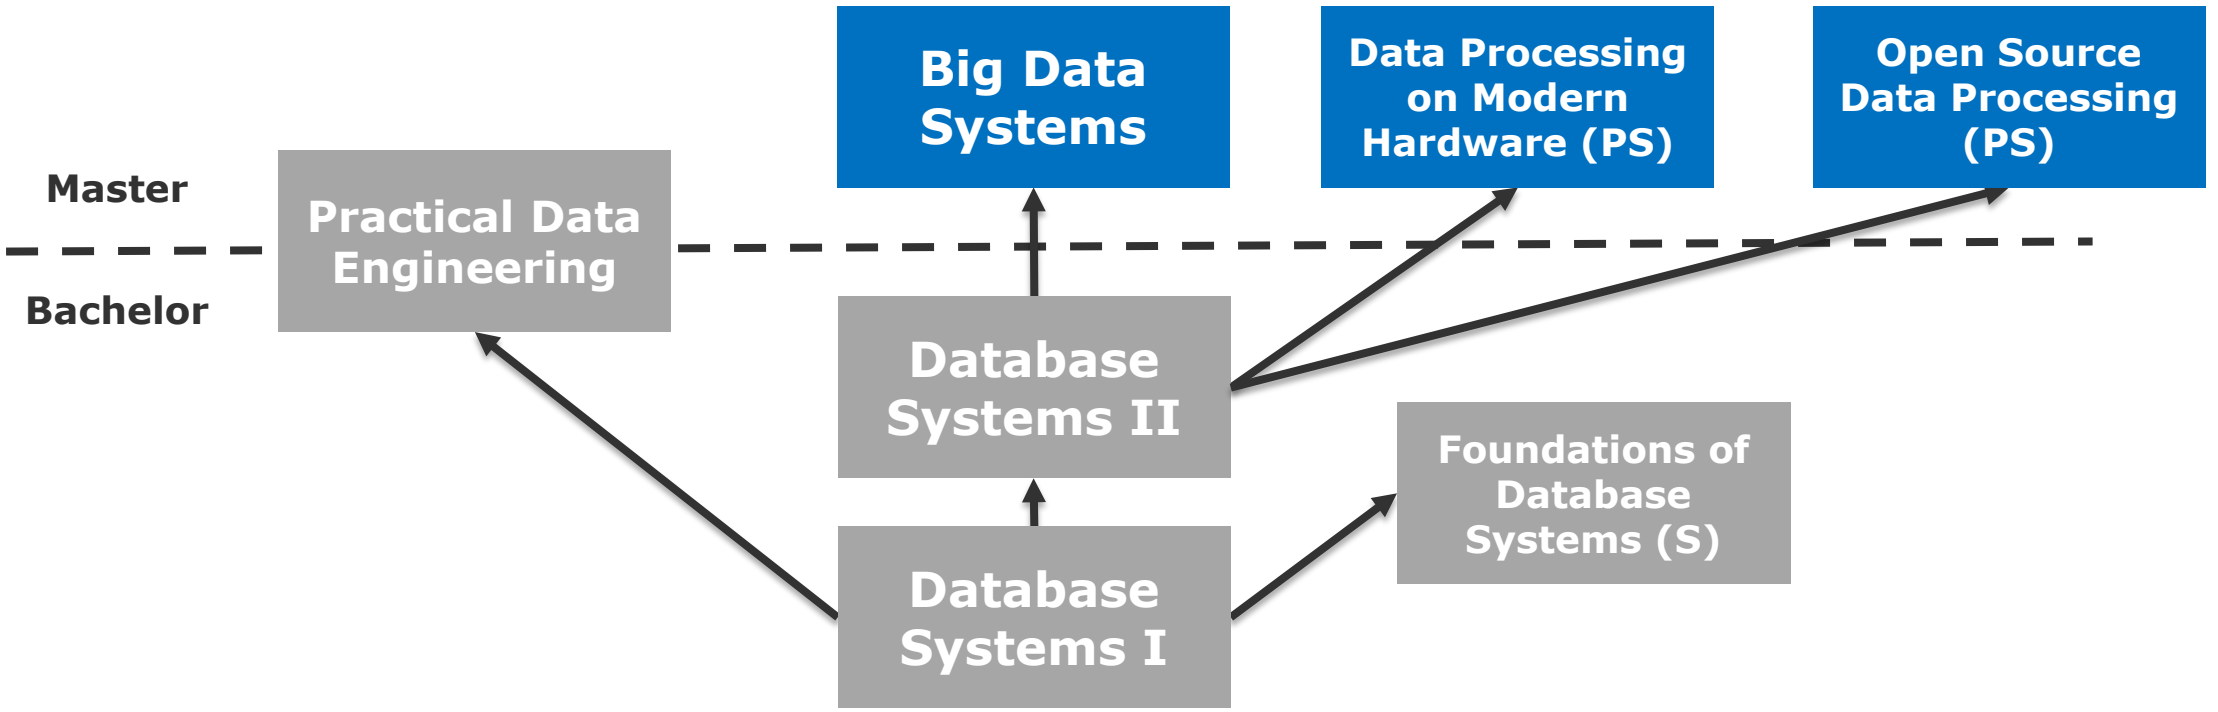

# Teaching

## This semester only Master level courses.

- Big Data Systems
  - Lecture on Implementation and Use of Big Data Systems
  - Tuesdays/Thursdays, 11:00 – 12:30
  - Master, 6 ECTS
- Open Source Data Processing
  - Learn how to do open source development with practitioners
  - Wednesdays, 11:00 – 12:30
  - Master, 6 ECTS
- B & M Projects
  - Bachelor Project – ML System Benchmarking
  - Master Project – Distributed Stream Processing
- Data Management on Modern Storage Technologies
  - Project Seminar on DB + PMem
  - Wednesdays, 13:30 – 15:00
  - Master, 6 ECTS
- GPU Accelerated Data Processing
  - Project Seminar on GPU + DP
  - Tuesdays, 13:30 – 15:00
  - Master, 6 ECTS
- Projects & Theses
  - Topics on Data Engineering Systems available
  - Guidance in Competitions and Challenges
  - SIGMOD SRC / Programming C / DEBS GC

# Motivation

# Big Fast Data

Data is growing

Messages, tweets, social networks (statuses, check-ins, shared content), blogs, click streams, various logs, …

- *Facebook: > 1,5B active users, > 60B messages/day*
- *Twitter: > 300M active users, > 500M tweets/day*

Everyone is interested!

**The value of data is decreasing with its age!**

# Stream Processing

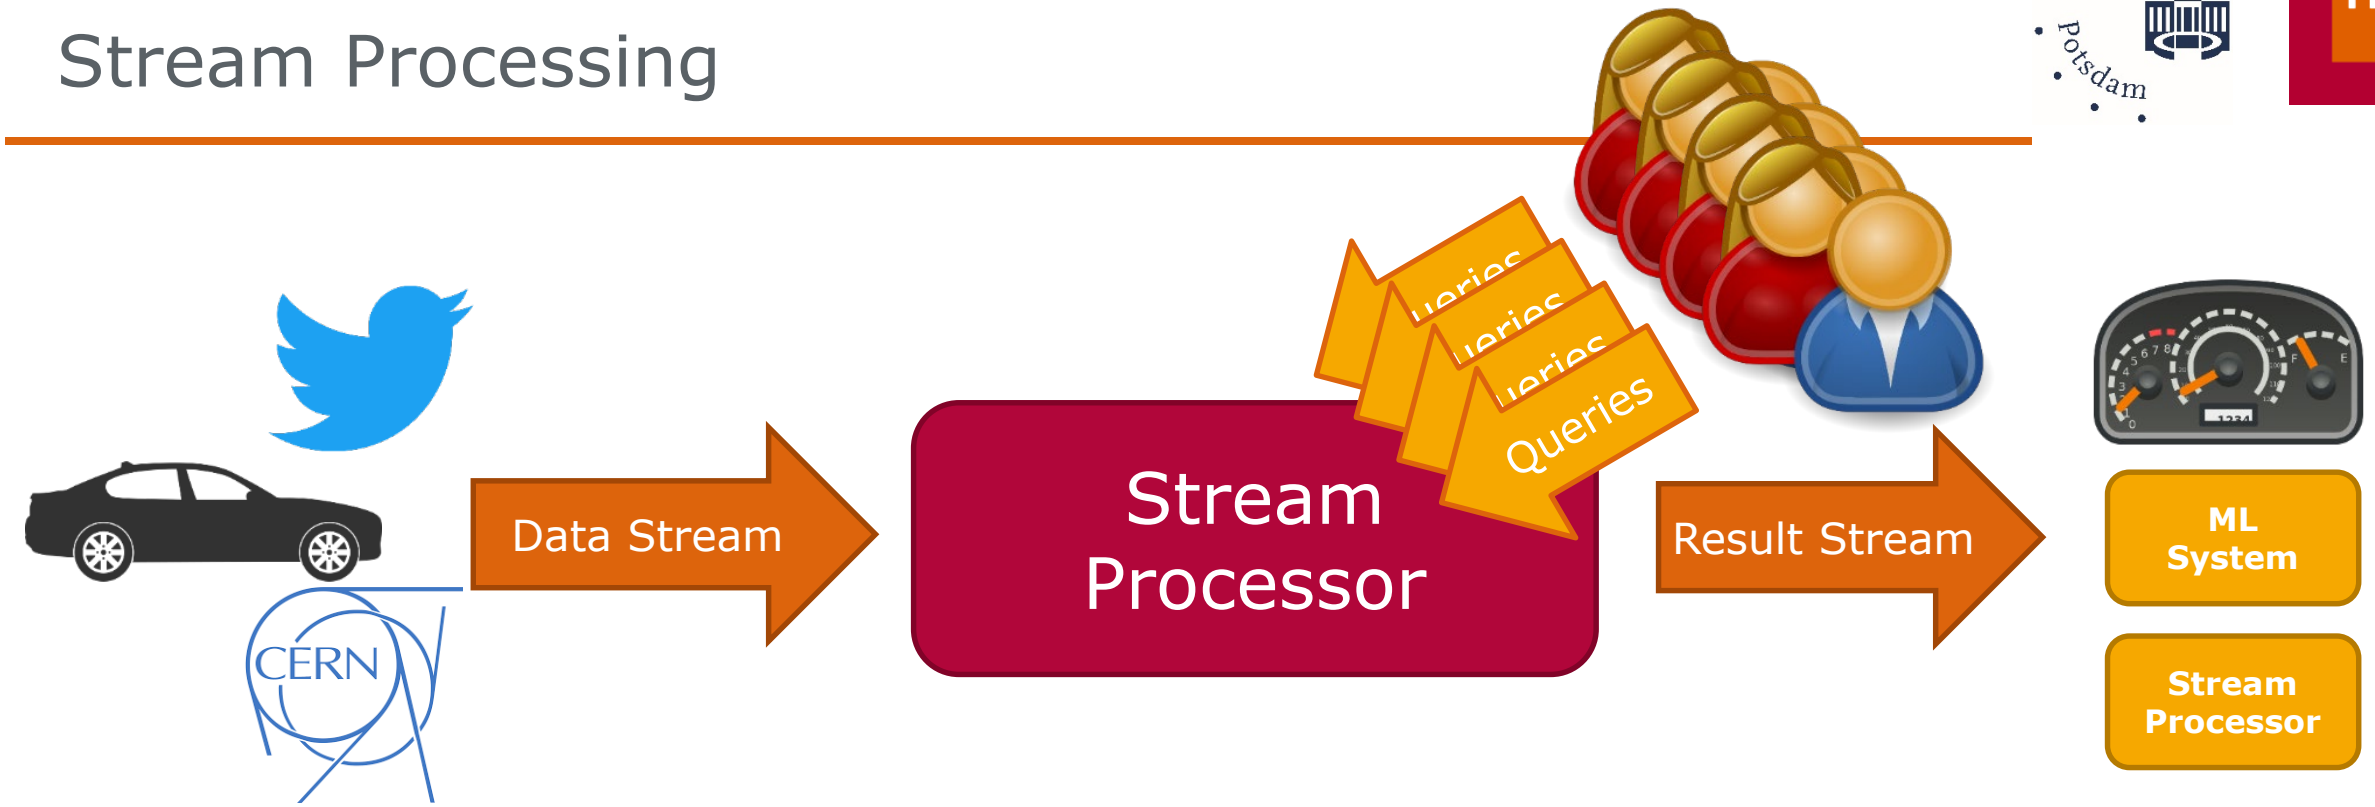


Challenge

- Potentially unlimited data set
- Many different queries
- Continuous results

# Streaming Processing Job

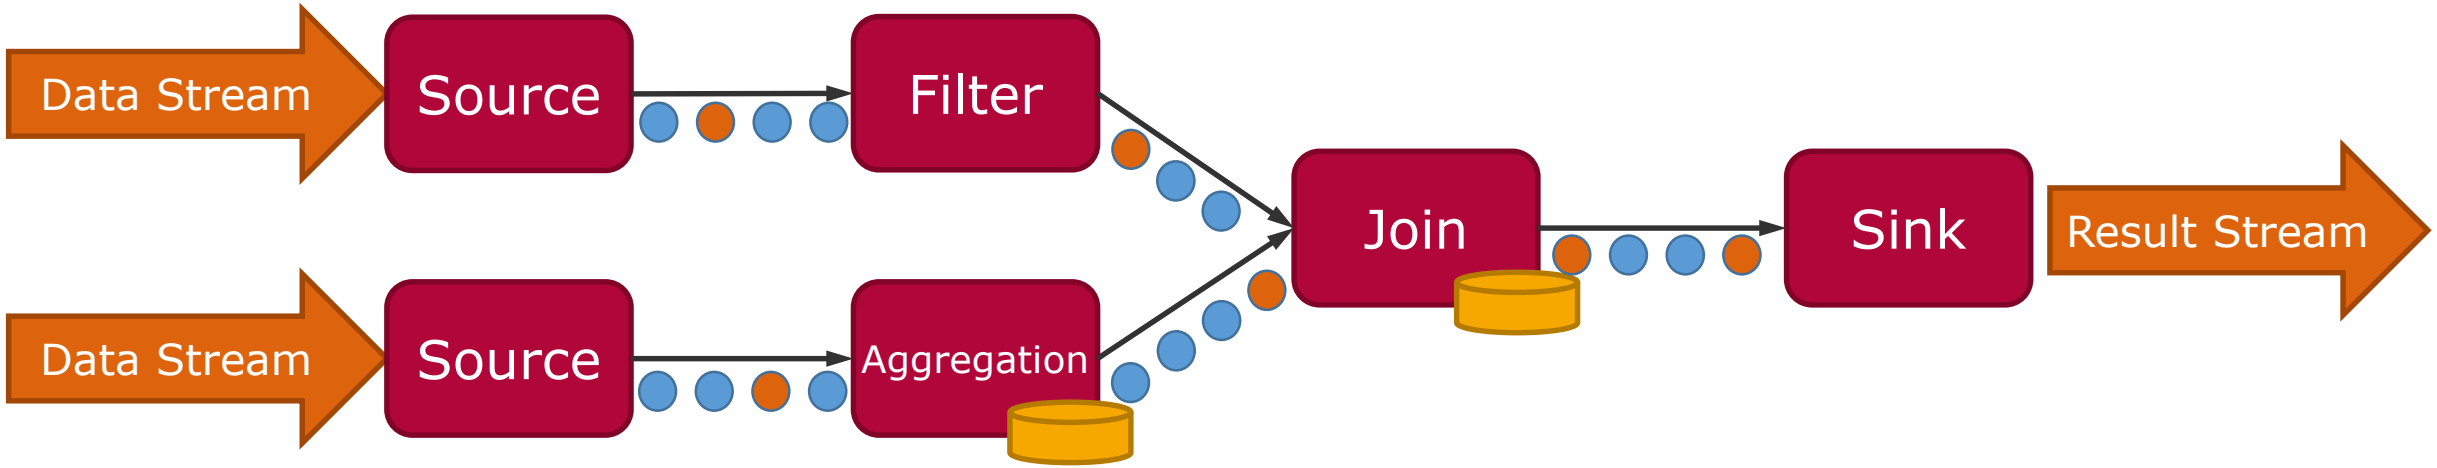


Dataflow

- Operators
- Records
- Control events
- State

# Distributed Stream Processing

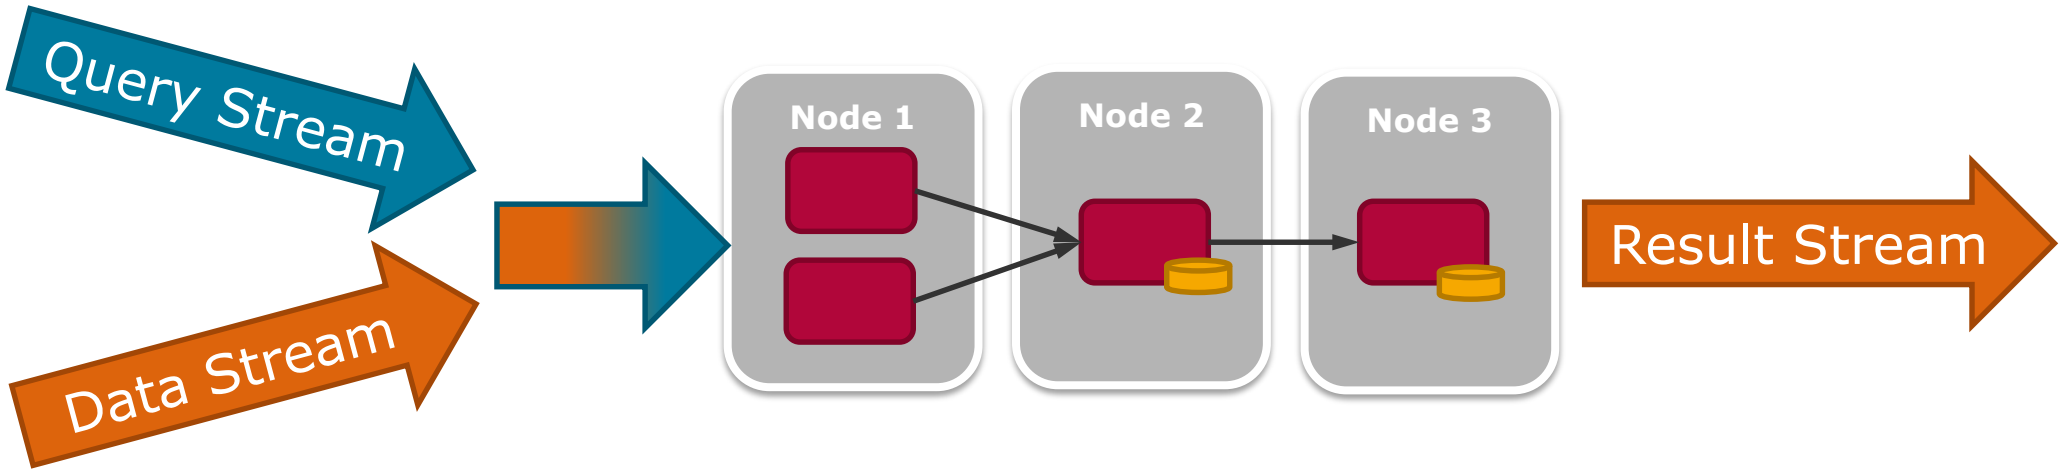


- Distribute query topology on multiple nodes
- Employ efficient network communication

Hendrik Makait: Rethinking Message Brokers on RDMA and NVM. SIGMOD 2020.

# Make it fast!

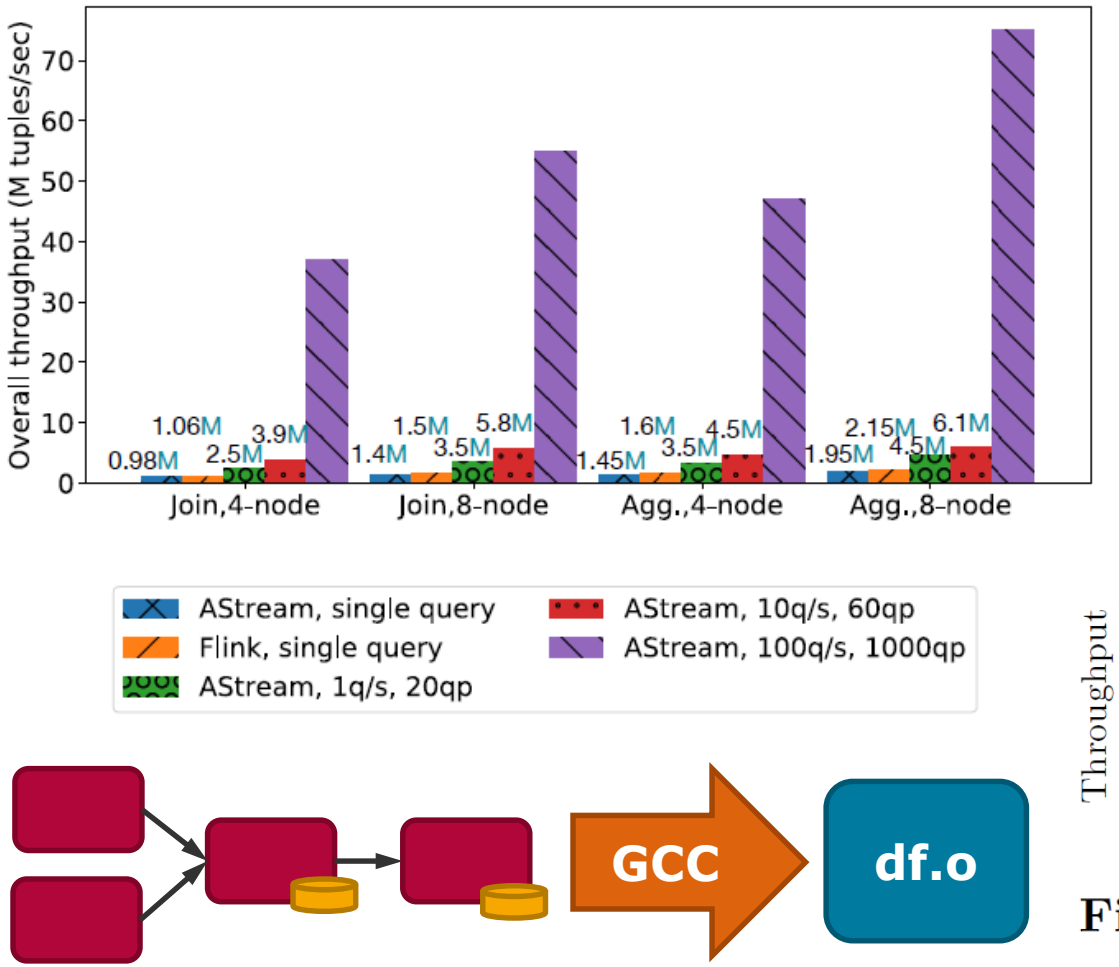





- Adhoc query processing is much more efficient!
- But it could be even faster – through code generation!
- Modern networking can help

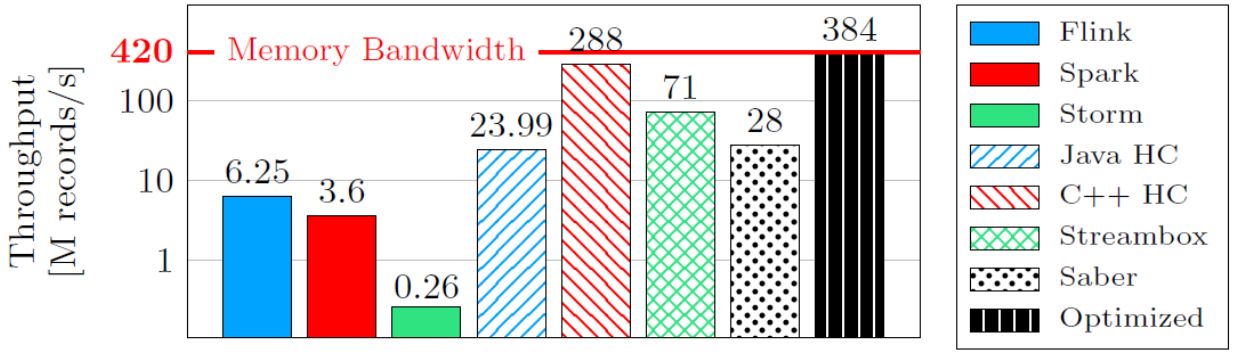


**Figure 1: Yahoo! Streaming Benchmark (1 Node).**

# Thank you!

Distributed Dynamic Stream Processing

Build

- Distributed stream processing engine prototype
- Query compilation pipeline
- ~~Multiquery optimization~~

Learn

- Distributed systems engineering
- Stream processing
- Query compilation

Supervising Team

- Pedro Silva
- Tilmann Rabl

# Meeting Time

- Proposal:
  - Wednesday 3:15 pm

# Code of Conduct

- This course should be fun for everyone
  - Needs to be a safe environment

- Communication
  - Learn how to write professional emails: https://medium.com/@lportwoodstacer/how-to-email-your-professor-without-being-annoying-af-cf64ae0e4087
  - Use netiquette in forum, email, chats, etc.

- Generally
  - Treat everyone with respect and consideration
  - HPI should be a safe place for everyone

# Project Overview and Milestones

# Overview

# Milestone 0: Work Environment (1 week)

- Main interests / expectations of each member
- Start reading the recommended literature
- Outline of group dynamics: e.g. hours per week, workdays of the week, etc.
- Weekly meeting time slot
- Software engineering approach: e.g. XP, Scrum, Waterfall, etc.
- Communication tools of the group: e.g. Slack, Discord, IRC, etc.
- IDEs: C Lion, Visual Studio, Eclipse, Emacs, etc.
- Repository solution: Github, Gitlab, Bitbucket, etc.
- File sharing: Google drive, Owncloud, Slack, etc.
- Quick presentation on the weekly meeting

# Milestone 1: Setting up the baseline (1-2 weeks)

- Defining an **use case/workload**: what the data will look like?
- Defining a **simple workflow on Flink** using only filter, join and window (aggregation) operators;
- Getting the code from Jeyhun's Flink benchmarking tool: https://github.com/streamline-eu/StreamBenchmarks
- Defining which hat **metrics** could be interesting to understand the performance of the execution.
- Running Jeyhun's benchmark on the application and **plotting the data**.
- **Short report and quick presentation** on the weekly meeting.

# Milestone 2: Proto 1, a first prototype (2-3 weeks)

- Defining a **hard-coded** "query compilation" for the workflow defined in Milestone 1 and implementing it.
- Defining **quick hardcoded procedures** to get the metrics from the query-compiled application
- Getting the metrics and **plotting the data** (using quick hardcoded scripts)
- **Comparing the results of benchmarking** the baseline (Milestone 1) and Proto 1
- Short report and quick presentation on the weekly meeting

# Milestone 3: Proto 2, more generic Proto 1 (4 weeks)

- Defining an approach for making the **workflow definition** more generic and **having parallel queries**
- Working on the **insights from Proto 1 benchmark results**: are there bottlenecks, or improvement opportunities? How hard it would be to implement them?
- Checking current considered metrics and planning and **defining new metrics** to help evaluating the impact of parallel queries and genericity
- **Improving the procedures** to get the metrics from Proto 2. Making them more reusable?
- **Improving the procedures** to plot the data of Proto 2's benchmark. Making them more reusable?
- **Comparing the results from benchmark** of Milestone 1 to Flink's and Proto 1's. How making it more generic affected the metrics?
- Short report and presentation on the weekly meeting

# Milestone 4: Proto 3, going distributed (4 weeks):

- **Analyzing the challenges** of distributing Proto 2 and defining a plan of attack
- **Implementing** the plan
- Checking current considered metrics and **defining new metrics** to help evaluating the impact of distribution on performance
- **Benchmarking** and comparing results to Flink's and Proto 2's
- Short report and presentation

# Bonus Milestone: Proto 4, the modern hardware one

- **Adding RDMA to the soup**. What are the benefits that RDMA could bring to Proto 3? And what are the challenges?
- **Designing** a plan of attack
- **Implementing** the plan
- **Benchmarking** and comparing results to Flink's, Proto 2's and Proto 3's
- Short report and presentation

# Milestone 5: Reporting in, team

- Time to wrap up and write the final report!

# An idea of the daring overview

# Thank you for your attention!

- Questions?
  - pedro.silva@hpi.de
  - tilmann.rabl@hpi.de